



CHINA LINUX KERNEL
中国Linux内核开发者大会



华中科技大学
网络安全学院
School of Cyber Science and Engineering, HUST

第19届中国 Linux内核开发者大会



赞助单位



支持单位



支持社区&媒体



2024年10月 湖北·武汉

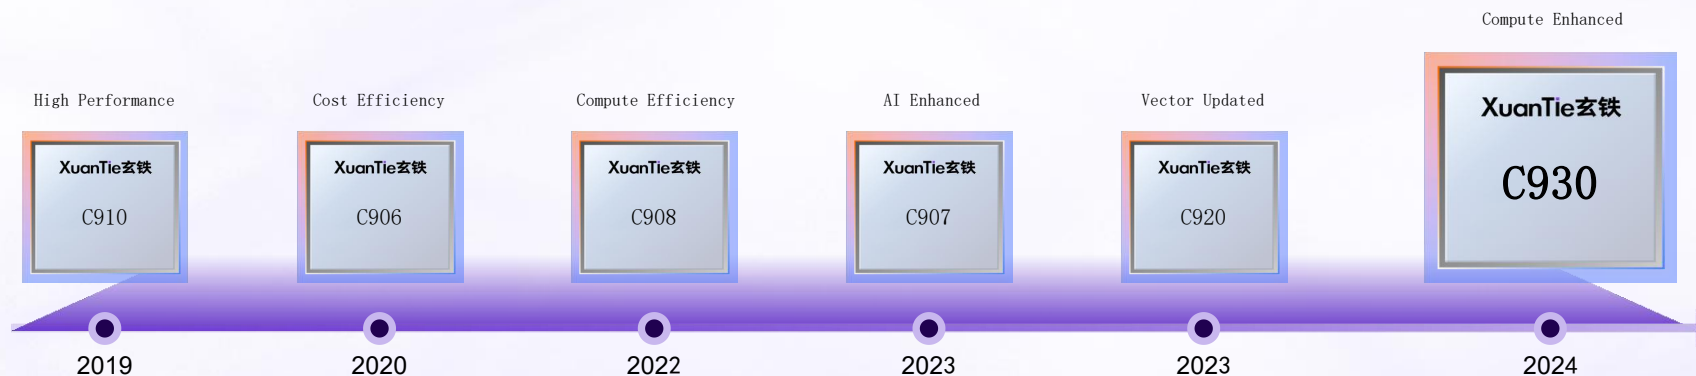


华中科技大学



玄铁是阿里巴巴达摩院旗下品牌，致力于推动RISC-V架构的前沿研究和建立在该架构基础上的开源生态的建立和发展，是国际 RISC-V 生态引领者，为数字化时代提供强大、智能、安全、开放的新型计算架构。

自成立以来，团队持续深耕 RISC-V 技术研发及生态建设，并陆续推出了一系列玄铁处理器，可满足高中低全系列性能需求。玄铁积极拥抱开源，坚持开放创新，已逐渐构建起以 RISC-V 为核心的生态体系，与生态伙伴协同推动 RISC-V 芯片、开发工具、操作系统、应用解决方案等不同层面的软硬一体化发展，全力推进 RISC-V 软硬全栈技术多领域发展落地，加速实现智能时代的万物互联！



RISC-V IOMMU 的 Scalable 提案 (GIPC) 及 Linux 技术路线

在 IO 虚拟化场景中，处理 PASID 的不同技术流派

郭任

阿里巴巴达摩院 高级技术专家



目录

Contents

01

什么是 PASID?

PCI-e TLP - 从 RID 到 PASID

02

Scalable Mode 提案

GIPC - Guest page table In Process Context

03

Linux 技术路线

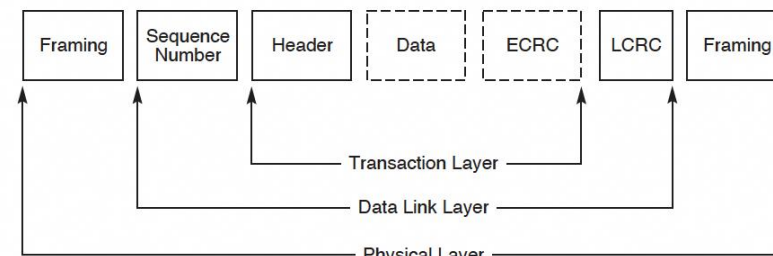

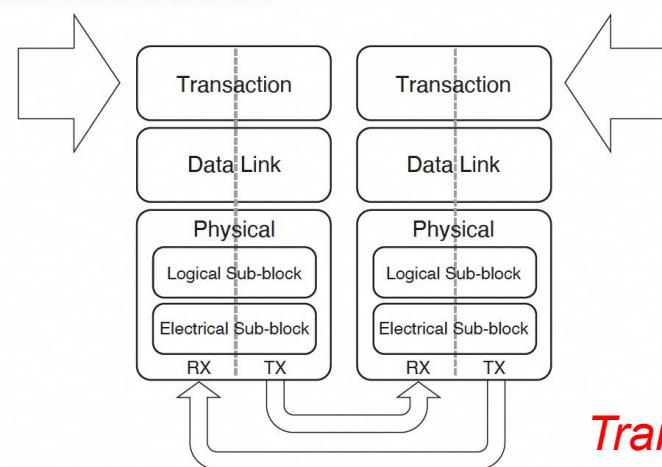
RISC-V IOMMU、emulated-IOMMU 和 virtio-IOMMU



01 | 什么是 PASID?

PCI-e TLP - 从 RID 到 PASID


从 PCI-e TLP 开始



Transaction Layer Packet




Address Space	Transaction Types	Basic Usage
Memory	Read Request/Completion, Write Request, Deferrable Write Request/Completion , AtomicOp Request/Completion	Transfer data to/from a memory-mapped location
I/O	Read Request/Completion, Write Request/Completion	Transfer data to/from an I/O-mapped location
Configuration	Read Request/Completion, Write Request/Completion	Device Function enumerating
Message	Baseline (including Vendor-Defined)	MSI IRQ




- **TLP Prefixes** 可选
- **TLP Header** 必选
- Data Payload 可选
- ECRC/Digest 可选

Header: Routing and Addressing Rules

ID routing is used with **Configuration Requests**, with **ID Routed Messages**, and with **Completions**.

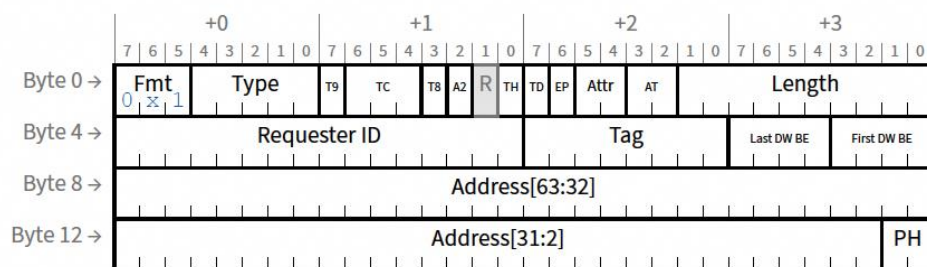


Header Format for Configuration Transactions

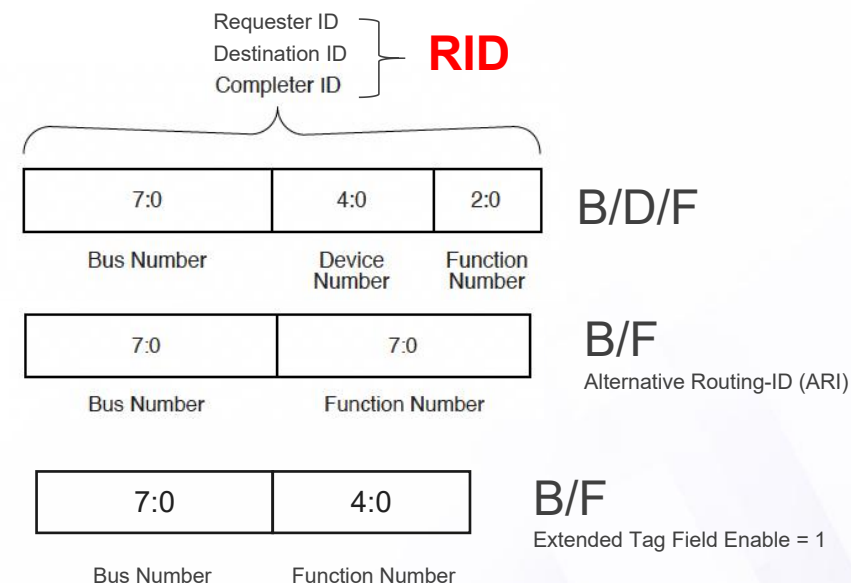


Header Format for Completion Transactions

(ID +) Address routing is used with **Memory, I/O Requests**:



Header Format for Memory Request Transactions



- Function 是 PCI-e 基于 RID 路由与寻址方式的最小设备单位。
- RID[16]: 最多 256 条总线，每条总线最多 256 个 Functions。
- 当 Extended Tag Field Enable = 1 时，位宽缩减 RID[13]。
- Function 在总线枚举阶段就已确定，每个 PF/VF 必须占用配置空间。
- 基于上述 Function 枚举原则的 SR-IOV (2006) 虚拟设备的成本高，且数量受限。



因此，基于 TLP Requester-ID 的 Scalable 方案正在起草：DRAFT PCI-SIG ECR - Scalable I/O Virtualization (by ArchProto WG, AMD, Ampere, Arm, Broadcom, Intel, Microsoft, NVIDIA, RedHat/IBM) (不兼容 SR-IOV)

<https://members.pcisig.com/wg/PCIe-Protocol/document/21183>

PASID TLP Prefix (Non-Flit Mode)

PASID 结合 RID 用于识别请求的地址空间：

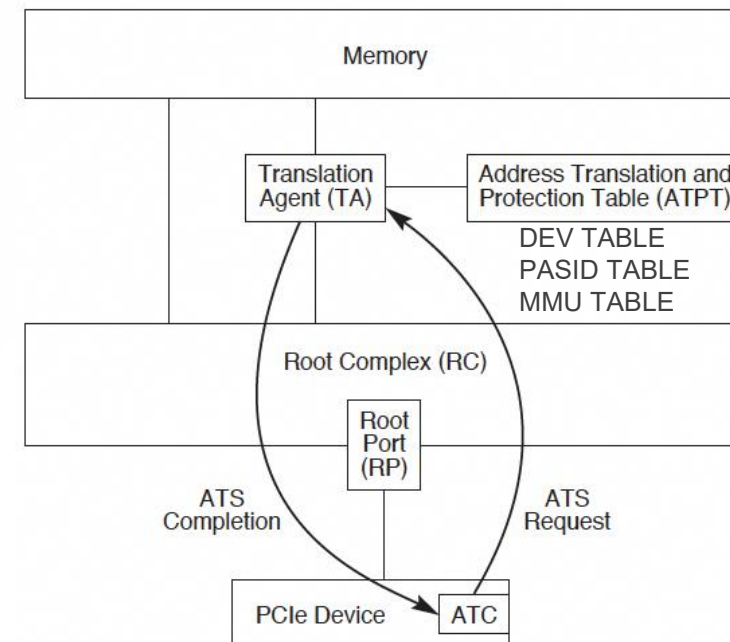


- 带有地址类型的 TLP 请求。
- ATS 请求 (MRd with AT=01b), ATS INV, ATS PRI 相关 Messages。



PASID TLP Prefix (Non-Flit Mode)

PASID 结合 RID 用于识别请求的地址空间：

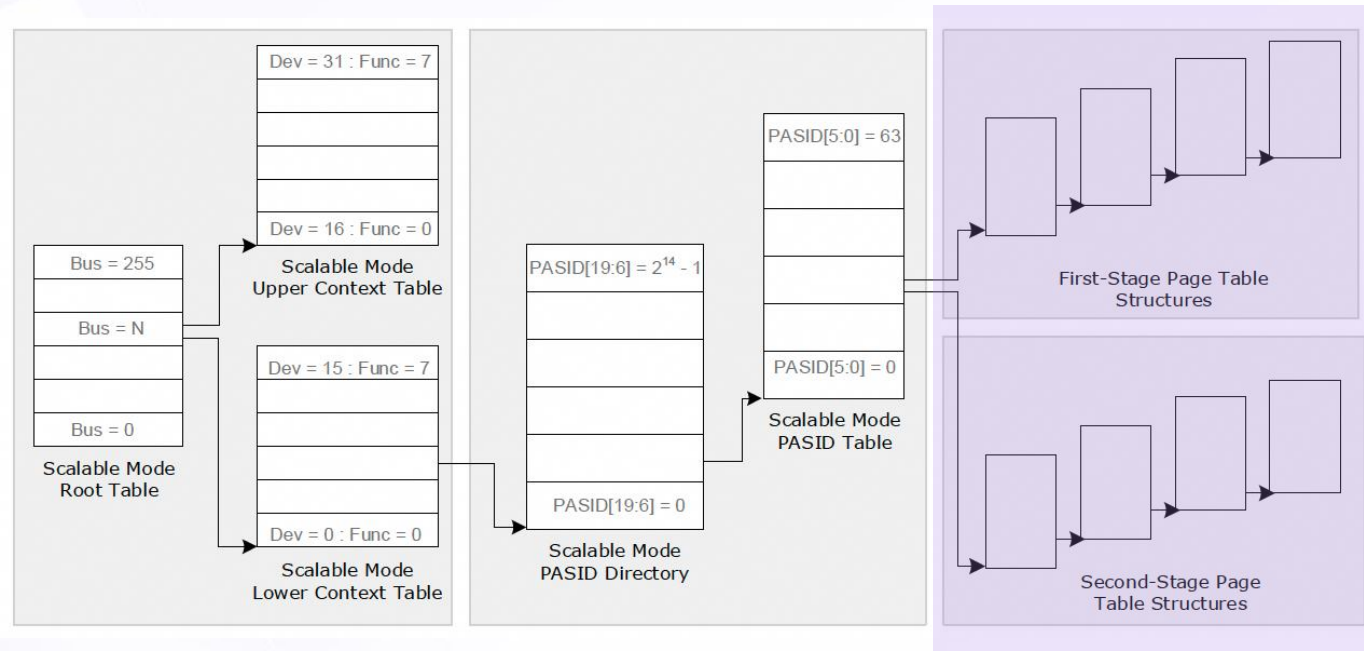
- 带有地址类型的 TLP 请求。
- ATS 请求 (MRd with AT=01b), ATS INV, ATS PRI 相关 Messages。



各家 TA 对 PASID 的处理形成两大技术流派：

1. Intel VT-d 的 Scalable Mode 方案
2. ARM-SMMU 和 AMD-Vi 的方案

When TA = Intel VT-d Scalable Mode




每一个 PASID Entry 拥有独立的 First-Stage Page Table 和 Second-Stage Page Table，拥有完整的两级地址空间。

Intel VT-d Scalable Mode Address Translation

June 2018, Intel VT-d Revision 3.0:

- Added support for scalable-mode translation for DMA Remapping, that enables PASID-granular first-level, second-level, nested and pass-through translation functions.

When TA = Intel VT-d Scalable Mode



Intel VT-d Scalable Mode Address Translation

June 2018, Intel VT-d Revision 3.0:

- Added support for scalable-mode translation for DMA Remapping, that enables PASID-granular first-level, second-level, nested and pass-through translation functions.



OPEN
Compute Project

Scalable I/O Virtualization Revision 1.0

Version 1.3

June 2022

Authors:

Intel Corporation (Contact Point: Tom Stachura)

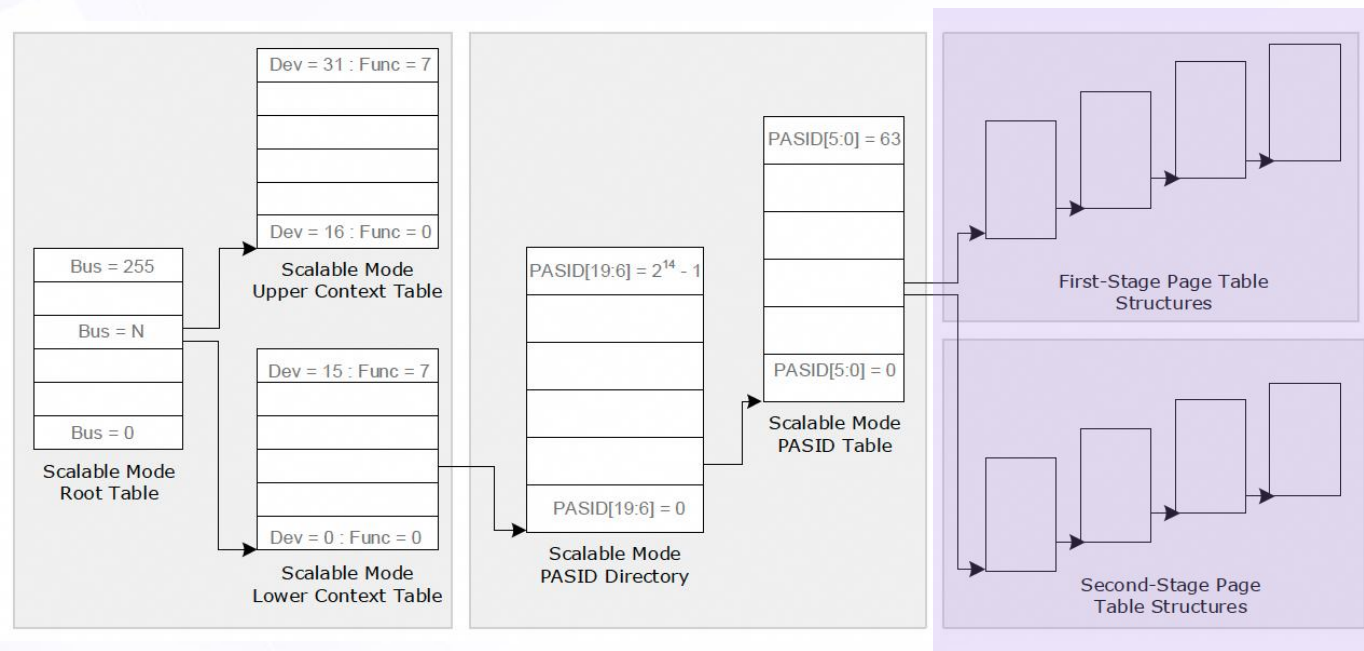
Microsoft Corporation (Contact Point: Renee L'Heureux)

June 2018: v1.0 by Intel

June 2022: v1.3 by OCP

PASID based ADI (Assignable Device Interfaces) 兼容 SR-IOV

When TA = Intel VT-d Scalable Mode

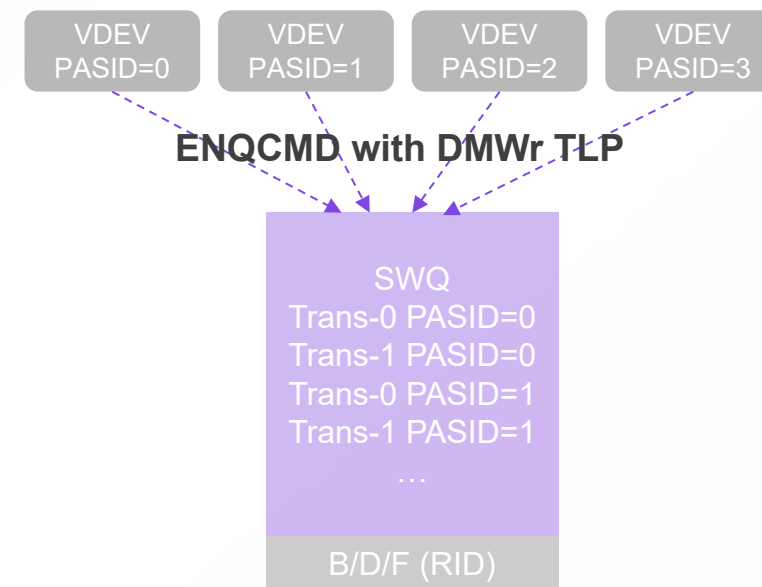


Intel VT-d Scalable Mode Address Translation

June 2018, Intel VT-d Revision 3.0:

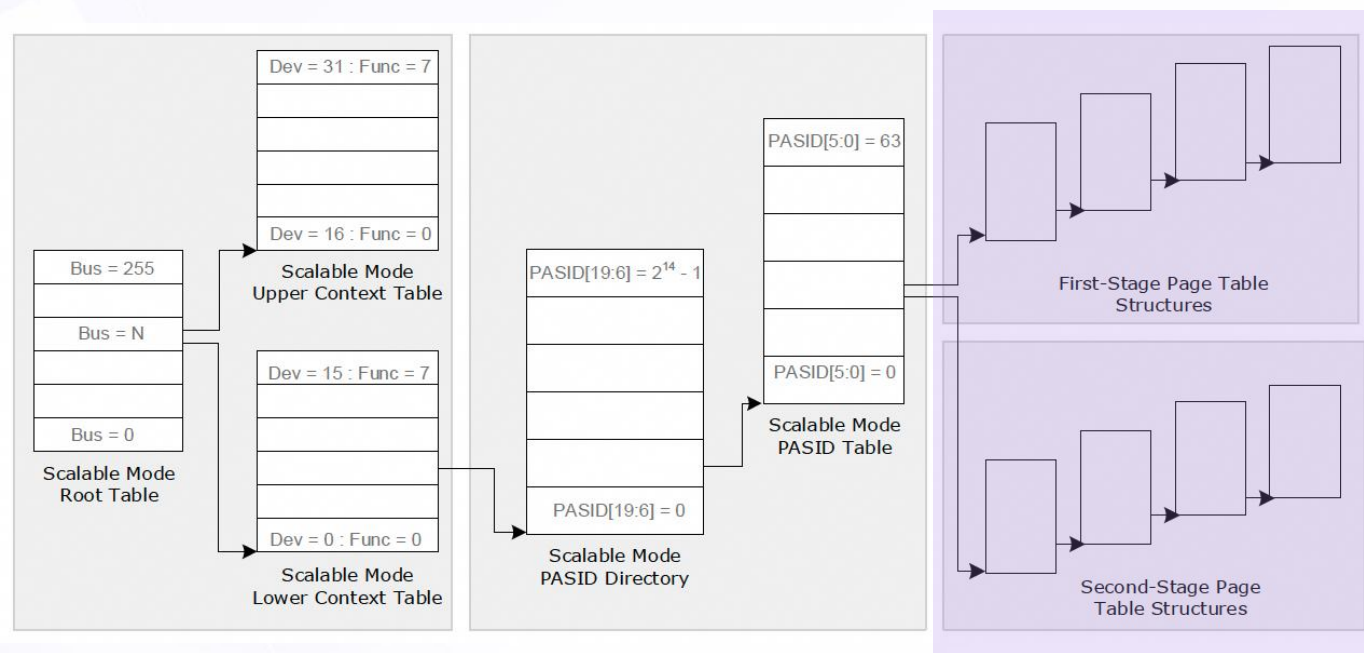
- Added support for scalable-mode translation for DMA Remapping, that enables PASID-granular first-level, second-level, nested and pass-through translation functions.

Scalable I/O Virtualization Spec: 7.3 ADIs Using Shared Work Queues



VT-d Scalable Mode 让同一个 Func 中的 VDEVs 可以自由分配给属于不同虚机的进程和容器。

When TA = Intel VT-d Scalable Mode

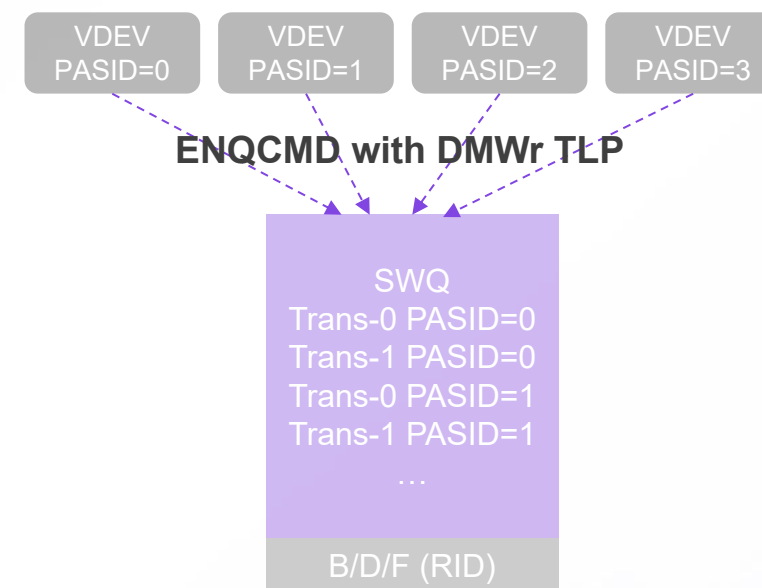


Intel VT-d Scalable Mode Address Translation

June 2018, Intel VT-d Revision 3.0:

- Added support for scalable-mode translation for DMA Remapping, that enables PASID-granular first-level, second-level, nested and pass-through translation functions.

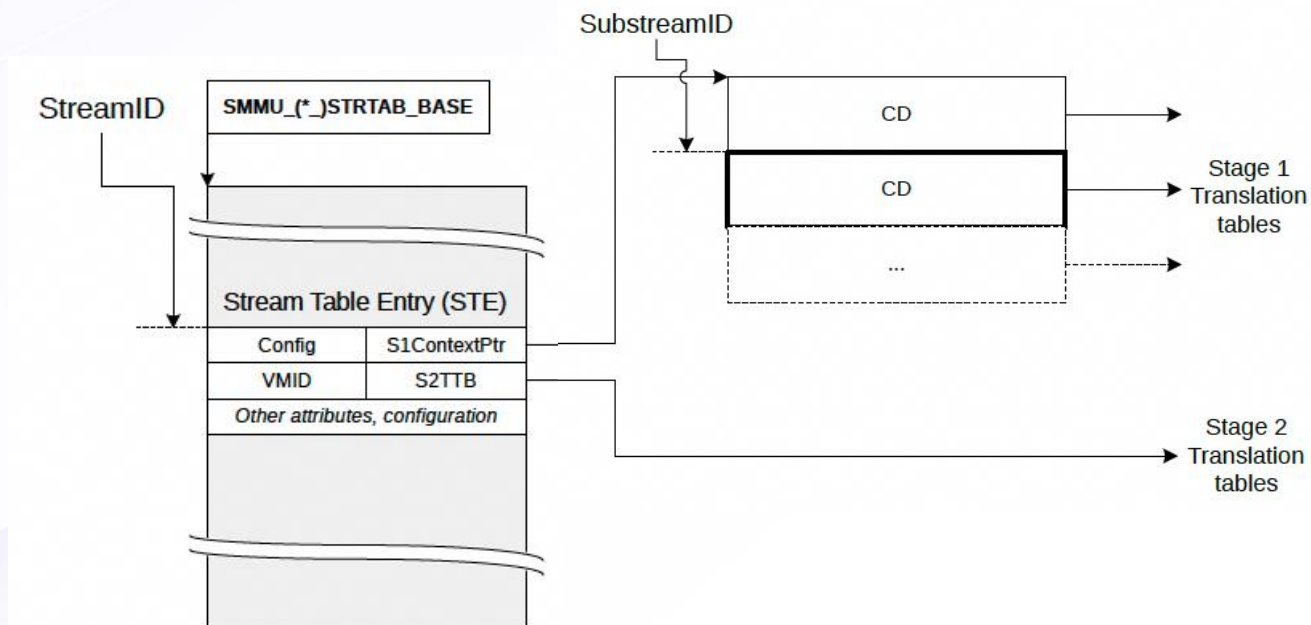
Scalable I/O Virtualization Spec:
7.3 ADIs Using Shared Work Queues



VT-d Scalable Mode 让同一个 Func 中的 VDEVs 可以自由分配给属于不同虚机的进程和容器。

对于同一个 Function 中的 VDEVs 进行跨虚机分配，ARM 和 AMD 有不同的见解 ...

When TA = ARM SMMUv3



StreamID Table -> PCI-e RID (B/D/F) index, Root Table
SubstreamID Table -> PCI-e PASID Directory

整张 SubstreamID (PASID) 表的 GPA-> PA 翻译, 被绑定在 STE.S2TTB 上, 因此**无法跨虚拟机分配 PASID**。




02 | Scalable Mode 提案 (RISC-V)

GIPC – Guest page table In Process Context

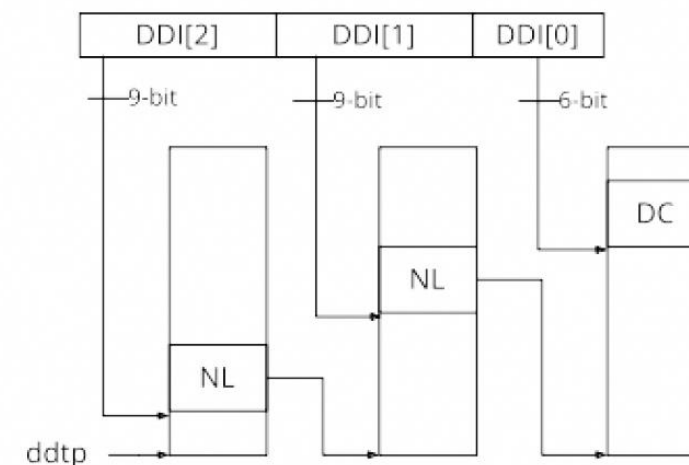
RISC-V IOMMU 表项设计

The PPN of the root device-directory-table is held in a memory-mapped register called the device-directory-table pointer (ddtp):

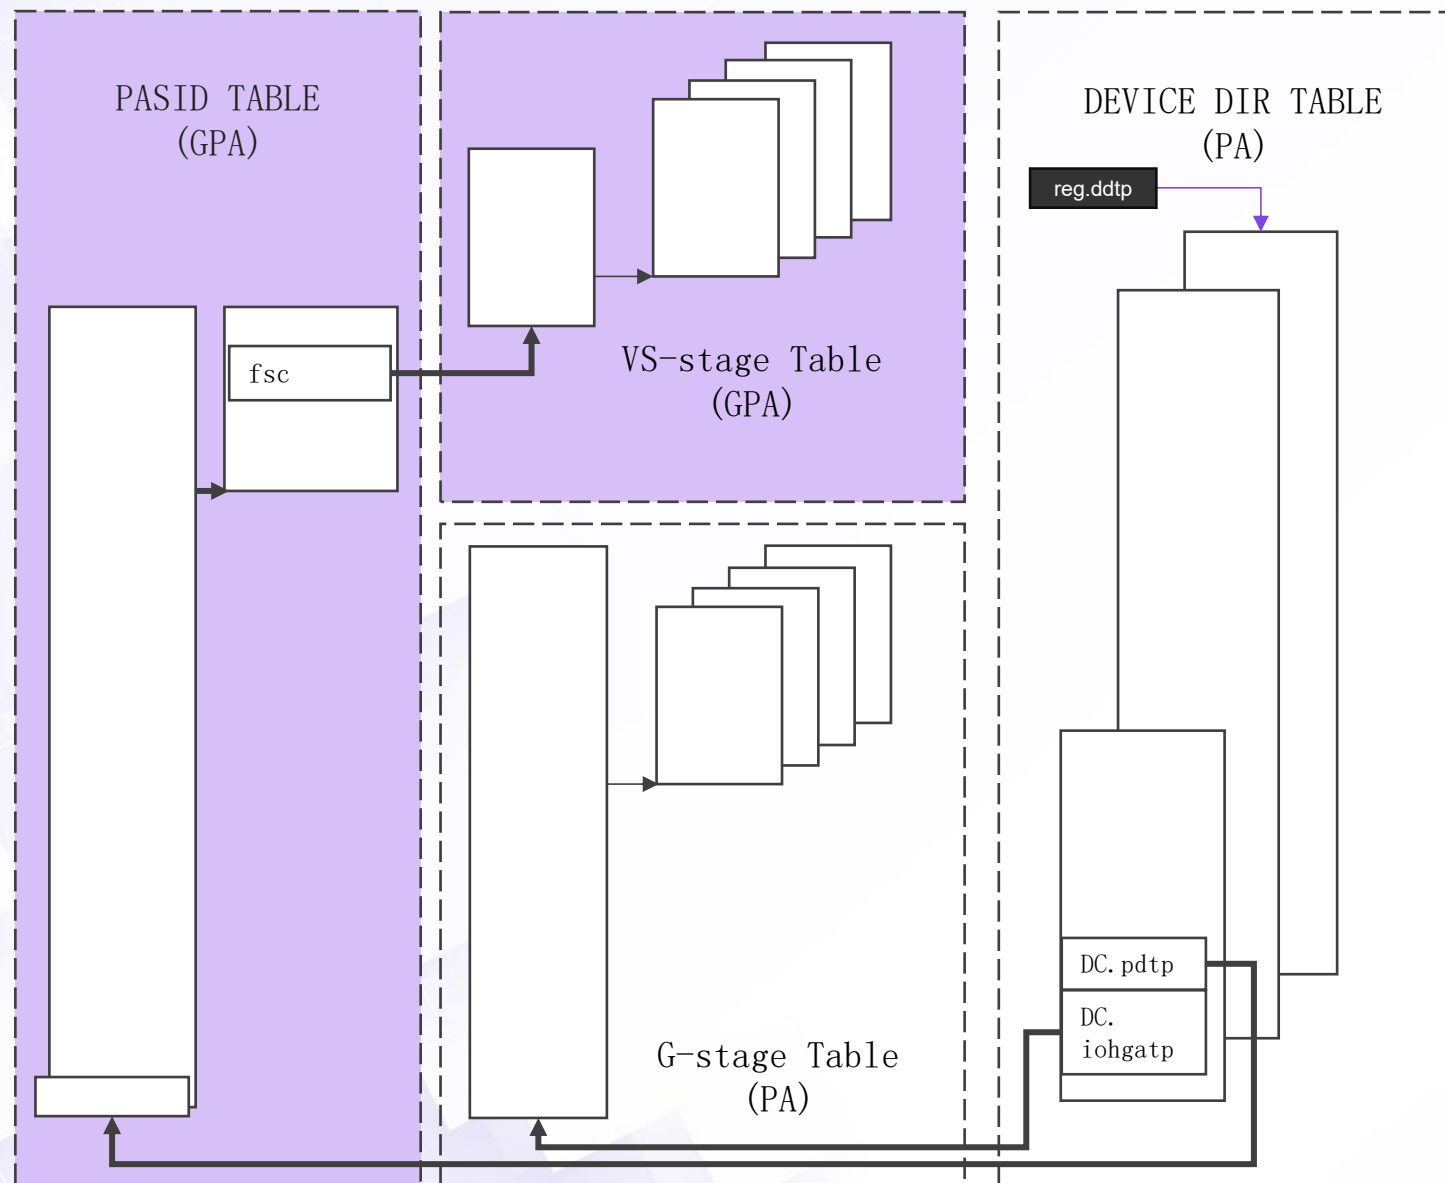


每一个 PCI-e Function, 对应 Device-Directory-Table (DDT) 中的一个 Device-Context (DC):

- Process-directory table pointer (pdtp) - PASID Table
- IO hypervisor guest address translation and protection (iohgatp) - Guest Page Table (GPA -> PA)



Three-level device directory with extended format DC.




类似 ARM-SMMU & AMD-Vi:

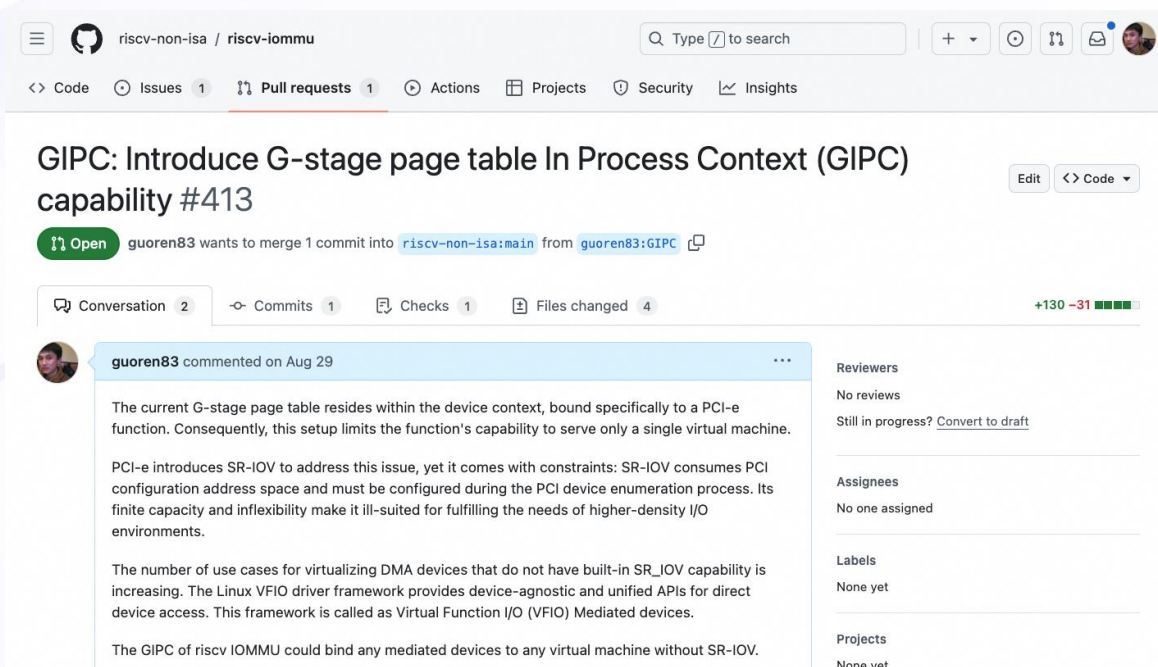
- 客户机管理整张 PASID 表
- PASID 表基于 GPA 寻址
- IOMMU 遍历 PASID 表需要类似 VS-stage 页表的二级地址翻译。

GIDC - Guest page table In Device Context

GIPC - Guest page table In Process Context



- 达摩院玄铁团队参与 RISC-V IOMMU Architecture Specification 初版制定，被列入贡献者名单。
- 如今，为满足高密 IO 场景需求，再次提出 Scalable 模式，以完善 RISC-V IOMMU 功能。



<https://github.com/riscv-non-isa/riscv-iommu/pull/413>

Authors:

Hao Ziyi <haoziyi@zhcomputing.com>

Guo Ren <guoren@linux.alibaba.com>

Song Zhuo <zhuo.song@linux.alibaba.com>

Dust Li <dust.li@linux.alibaba.com>

Shuai Xue <xueshuai@linux.alibaba.com>

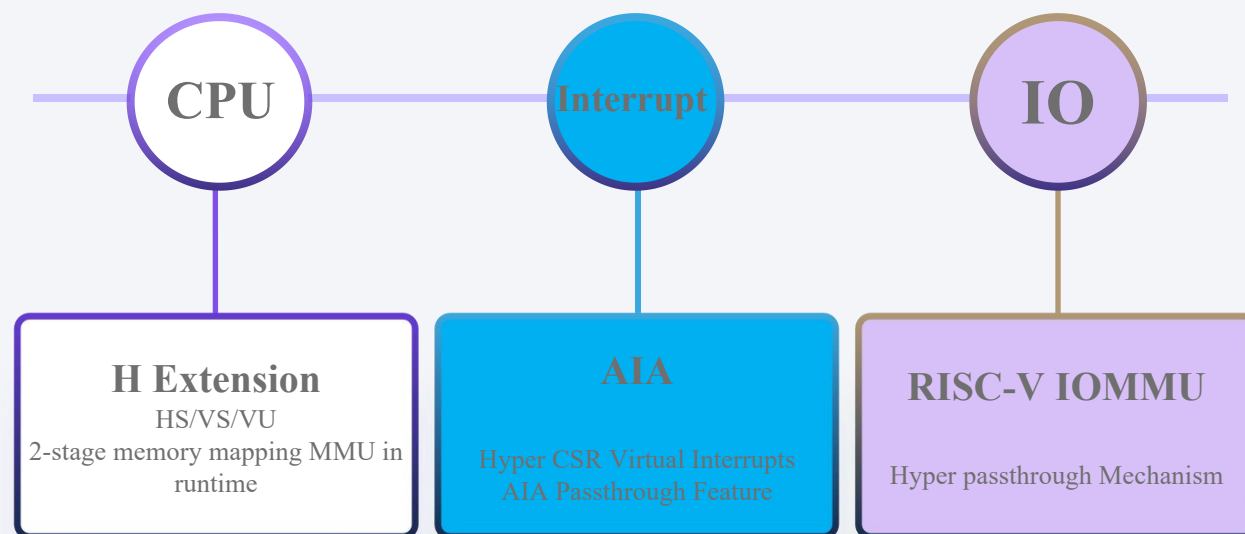
Feng Guanghui

<guanghuifeng@linux.alibaba.com>

Kaige Fu <kaige.fu@linux.alibaba.com>

Hao Xiang <hao.xiang@linux.alibaba.com>

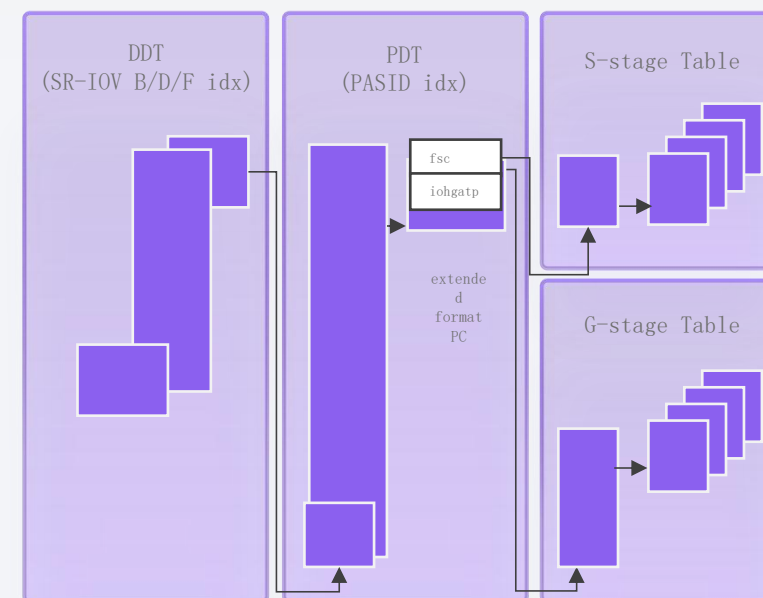
XuanTie RISC-V Virtualization Solution



vCPU Performance Overhead
from core binding < 1%

Interrupt Forwarding Overhead
in Virtual Machines ~ 0%

DMA Performance Overhead < 5%



Scalable IOMMU - Extension to Meet the
Demands of High-Density I/O Scenarios

- Implementation of Scalable IOMMU
- IRQ-Remapping
- **RISC-V virtio-IOMMU**




03

Linux 技术路线

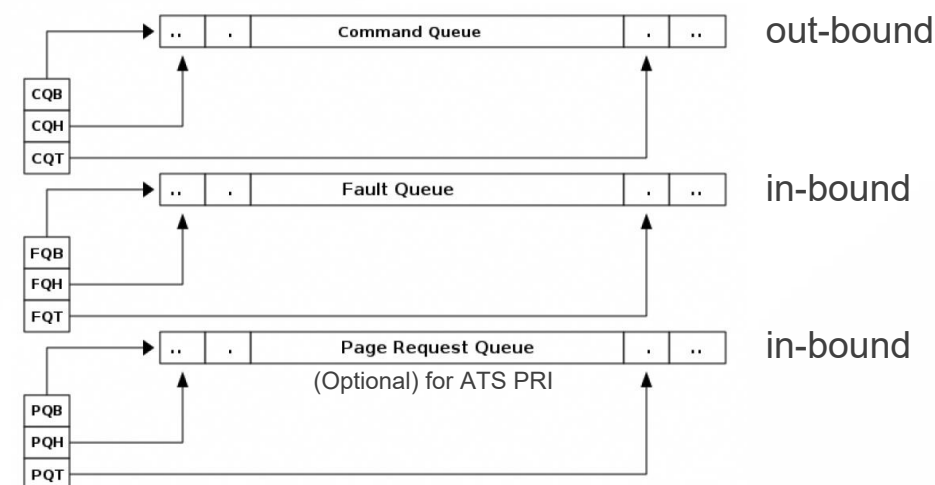
RISC-V IOMMU、emulated-IOMMU 和 virtio-IOMMU

RISC-V IOMMU & Linux Driver

Offset	Name	Size	Description	Is Optional?
0	capabilities	8	Capabilities of the IOMMU	No
8	ftl	4	Features control	No
12	custom	4	Designated For custom use	
16	ddtp	8	Device directory table pointer	No
24	cqb	8	Command-queue base	No
32	cqh	4	Command-queue head	No
36	cqt	4	Command-queue tail	No
40	fqb	8	Fault-queue base	No
48	fqh	4	Fault-queue head	No
52	fqt	4	Fault-queue tail	No
56	pqb	8	Page-request-queue base	if capabilities.ATS==0
64	pqh	4	Page-request-queue head	if capabilities.ATS==0
68	pqt	4	Page-request-queue tail	if capabilities.ATS==0
72	cqcsr	4	Command-queue CSR	No
76	fqcsr	4	Fault-queue CSR	No
80	pqcsr	4	Page-request-queue CSR	if capabilities.ATS==0
84	ipsr	4	Interrupt pending status register	No
88	iocntovf	4	HPM counter overflows	if capabilities.HPM==0
92	iocntinh	4	HPM counter inhibits	if capabilities.HPM==0
96	iohpmcycles	8	HPM cycles counter	if capabilities.HPM==0
104	iohpmctr1-31	248	HPM event counters	if capabilities.HPM==0
352	iohpmevt1-31	248	HPM event selector	if capabilities.HPM==0
600	tr_req_iova	8	Translation-request IOVA	if capabilities.DBG==0
608	tr_req_ctl	8	Translation-request control	if capabilities.DBG==0
616	tr_response	8	Translation-request response	if capabilities.DBG==0
624	Reserved	64	Reserved for future use (WPRI)	
688	custom	72	Designated for custom use (WARL)	
760	icvec	8	Interrupt cause to vector register	No
768	msi_cfg_tbl	256	MSI Configuration Table	if capabilities.IGS==WSI
1024	Reserved	3072	Reserved for standard use	



- Data Structures (DDT, DC, PDT, PC ...)
- In-memory queue interface (CQ, FQ, PQ)
- Memory-mapped register interface



[PATCH v10 0/7] Linux RISC-V IOMMU Support
 by Tomasz Jeznach <tjeznach@rivosinc.com>
<https://lore.kernel.org/linux-riscv/cover.1729059707.git.tjeznach@rivosinc.com/>

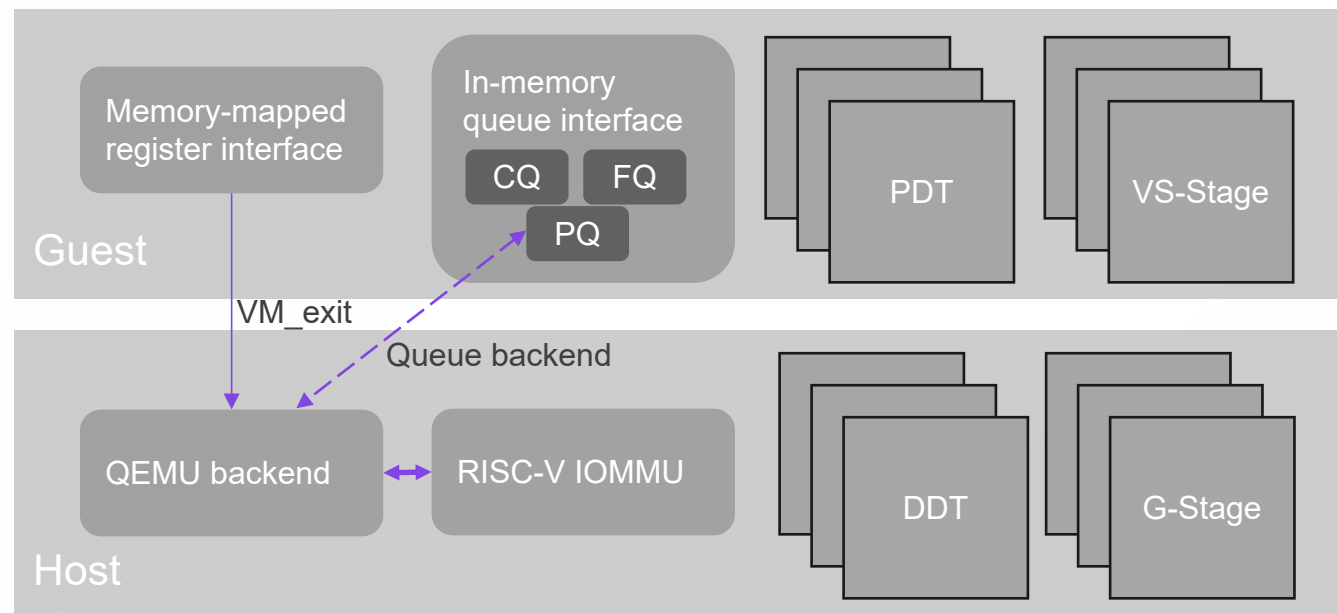
- Support platform device driver
- Support PCIe device driver

- Guest 对 Memory-mapped register 访问, 都会陷入 Hypervisor 处理
- QEMU 后端填充/读取 CQ, FQ, PQ。

[RFC PATCH v2 00/10] RISC-V IOMMU HPM and nested IOMMU support
by Zong Li <zong.li@sifive.com>
<https://lore.kernel.org/lkml/20240614142156.29420-1-zong.li@sifive.com/>

- Support GPA -> SPA (Tested)
- Support nested iommu (Not ready for Test)

但是, 尚未给出 QEMU backend 代码, 所以暂时无法测试 nested-translation。



缺点:

- 访问 Memory-mapped 寄存器效率低, 产生 VM_exit。
 - Guest 复用 RISC-V IOMMU 驱动, Nested 设计增加驱动复杂性。
 - 各个架构都要实现自己的 Qemu 后端驱动, 工作量大, 且不统一。
- 短期不会支持 GIPC, 我们认为有更好的选择:

Virtio-iommu specification

With page table extensions

06 October 2023

5.13.6.4 DETACH/ATTACH_TABLE request

```
struct virtio_iommu_req_attach_table {
    struct virtio_iommu_req_head head;
    le32 domain;
    le32 endpoint;
    u8 format;
    u8 reserved[111];
    struct virtio_iommu_req_tail tail;
};

#define VIRTIO_IOMMU_ATTACH_TABLE_ARM_SMMU3 1
#define VIRTIO_IOMMU_ATTACH_TABLE_INTEL_PT 2
#define VIRTIO_IOMMU_ATTACH_TABLE_RISCV 4
#define VIRTIO_IOMMU_ATTACH_TABLE_AMD_GCR3 5
#define VIRTIO_IOMMU_ATTACH_TABLE_AMD_PT 6
```

5.13.7	Acceleration	
5.13.7.1	Arm SMMUv3 tables	
5.13.7.1.1	PROBE property for Arm SMMUv3	
5.13.7.1.2	ATTACH_TABLE request for Arm SMMUv3	
5.13.7.1.3	INVALIDATE request for Arm SMMUv3	
5.13.7.2	Intel VT-d tables	
5.13.7.2.1	PROBE property for VT-d page tables	
5.13.7.2.2	ATTACH_TABLE request for VT-d page table	
5.13.7.2.3	INVALIDATE request for VT-d page table	
5.13.7.3	RISC-V IOMMU tables	
5.13.7.3.1	PROBE properties for RISC-V IOMMU	
5.13.7.3.2	ATTACH_TABLE request for RISC-V IOMMU	
5.13.7.3.3	INVALIDATE request for RISC-V IOMMU	
5.13.7.4	AMD GCR3 table	
5.13.7.4.1	PROBE properties for AMD GCR3 table	
5.13.7.4.2	ATTACH_TABLE request for AMD GCR3 table	
5.13.7.4.3	ATTACH_TABLE request for AMD page table	

5.13.7.3.2 ATTACH_TABLE request for RISC-V IOMMU

```
struct virtio_iommu_req_attach_table_riscv {
    struct virtio_iommu_req_head head;
    le32 domain;
    le32 endpoint;
    u8 format;
    u8 reserved[3];
    le64 tc;
    le64 ta;
    le64 fsc;
    u8 reserved[84];
    struct virtio_iommu_req_tail tail;
};

#define VIRTIO_IOMMU_HW_RISCV_TC_PDTV (1 << 5)
#define VIRTIO_IOMMU_HW_RISCV_TC_SADE (1 << 8)
#define VIRTIO_IOMMU_HW_RISCV_TC_DPE (1 << 9)
#define VIRTIO_IOMMU_HW_RISCV_TC_SBE (1 << 10)
#define VIRTIO_IOMMU_HW_RISCV_TC_SXL (1 << 11)

#define VIRTIO_IOMMU_HW_RISCV_TA_PSCID_SHIFT 12
#define VIRTIO_IOMMU_HW_RISCV_TA_PSCID_MASK 0xffff

#define VIRTIO_IOMMU_HW_RISCV_FSC_MODE_SHIFT 60
#define VIRTIO_IOMMU_HW_RISCV_FSC_MODE_MASK 0xf
#define VIRTIO_IOMMU_HW_RISCV_FSC_PPN_MASK 0xfffffffffff

/* When VIRTIO_IOMMU_HW_RISCV_TC_PDTV == 0 */
#define VIRTIO_IOMMU_HW_RISCV_FSC_MODE_SV39 8
#define VIRTIO_IOMMU_HW_RISCV_FSC_MODE_SV48 9
#define VIRTIO_IOMMU_HW_RISCV_FSC_MODE_SV57 10
#define VIRTIO_IOMMU_HW_RISCV_FSC_MODE_SV32 8

/* When VIRTIO_IOMMU_HW_RISCV_TC_PDTV == 1 */
#define VIRTIO_IOMMU_HW_RISCV_FSC_MODE_PD8 1
#define VIRTIO_IOMMU_HW_RISCV_FSC_MODE_PD17 2
#define VIRTIO_IOMMU_HW_RISCV_FSC_MODE_PD20 3
```

iommu: Add virtio-iommu driver (2019.1)
 Author: Jean-Philippe Brucker
jean-philippe@linaro.org

优势:

- Linux 五年前已支持 virtio-IOMMU
- QEMU 后端框架已实现
- 无需模拟 RISC-V IOMMU 实现
- virtio 命令队列更高效
- 跨 CPU ISA 支持
- 支持 GIPC 和 GIDC

- RISC-V IOMMU 仅支持 GIDC
- GIPC 提案使 RISC-V IOMMU 可以同时支持 GIPC 和 GIDC 的混合部署。
- 支持 Emulated-IOMMU
- 支持 Virtio-IOMMU（欢迎一起共建）

Thank you



玄铁公众号



玄铁中文站



玄铁海外站

