



CHINA LINUX KERNEL
中国Linux内核开发者大会



华中科技大学
网络安全学院
School of Cyber Science and Engineering, HUST

第19届中国 Linux内核开发者大会



赞助单位



支持单位



支持社区&媒体



2024年10月 湖北·武汉



华中科技大学

RISC-V eBPF 现状与未来

蒲乐辉

Pu Lehui <pulehui@huawei.com>





CHINA LINUX KERNEL
中国Linux内核开发者大会

目录

01 RISC-V eBPF 背景

02 RISC-V eBPF 现状

03 RISC-V eBPF 规划



CHINA LINUX KERNEL
中国Linux内核开发者大会

01 RISC-V eBPF 背景

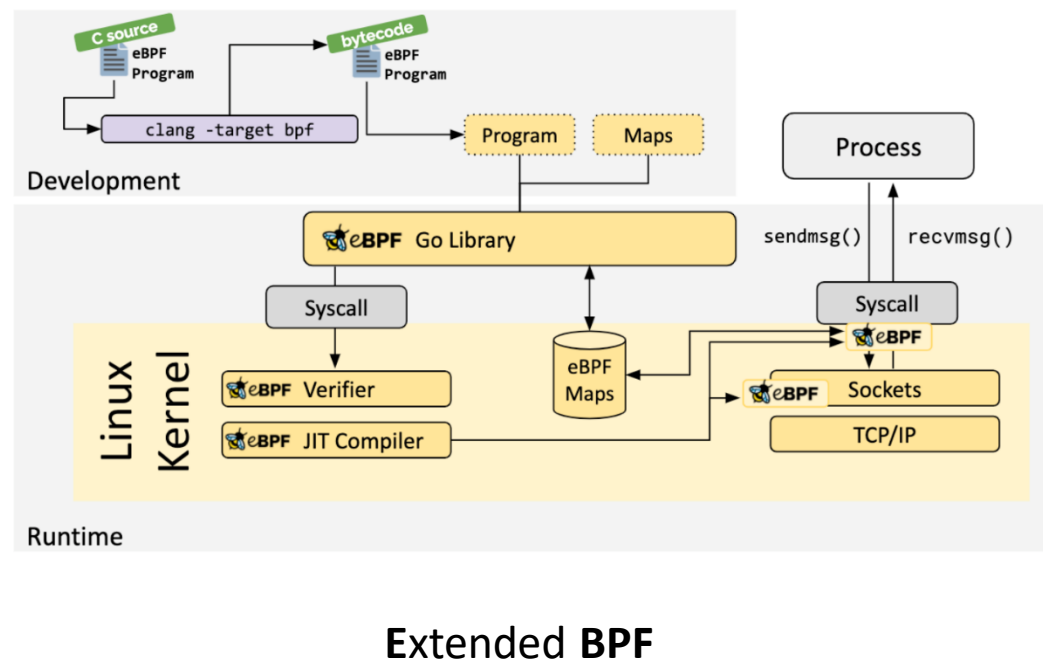
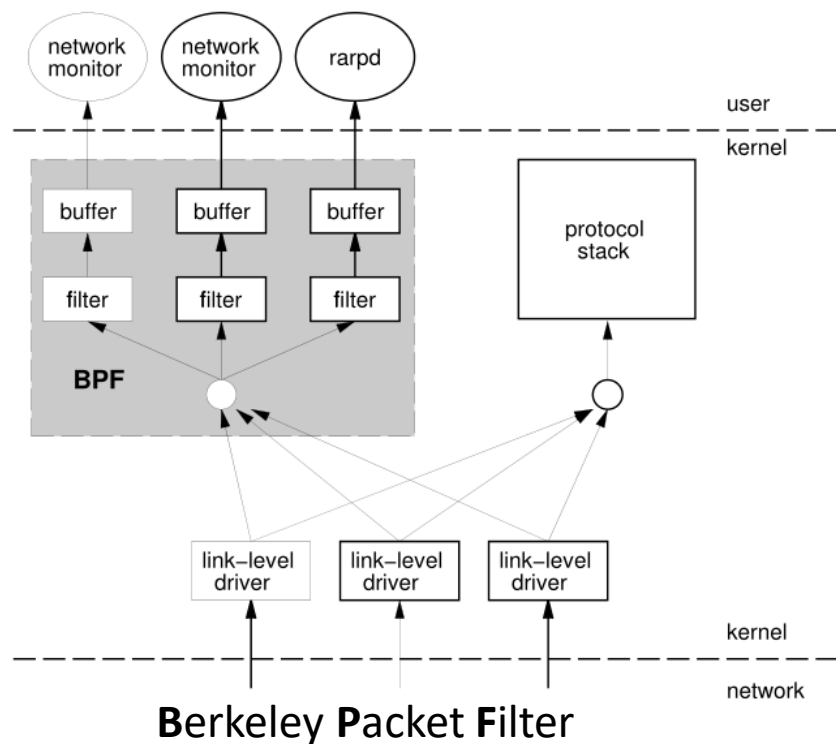
eBPF 介绍



CHINA LINUX KERNEL
中国Linux内核开发者大会

BPF是一种基于虚拟机模型的网络报文过滤器，通过将用户态定义的过滤规则编译成字节码注入内核，内核态虚拟机执行这段字节码完成报文过滤。

eBPF衍生自网络报文过滤器BPF。通过扩充寄存器数量、位宽，引入verifier和MAP机制，将BPF扩展为一个**通用**的内核态虚拟机。

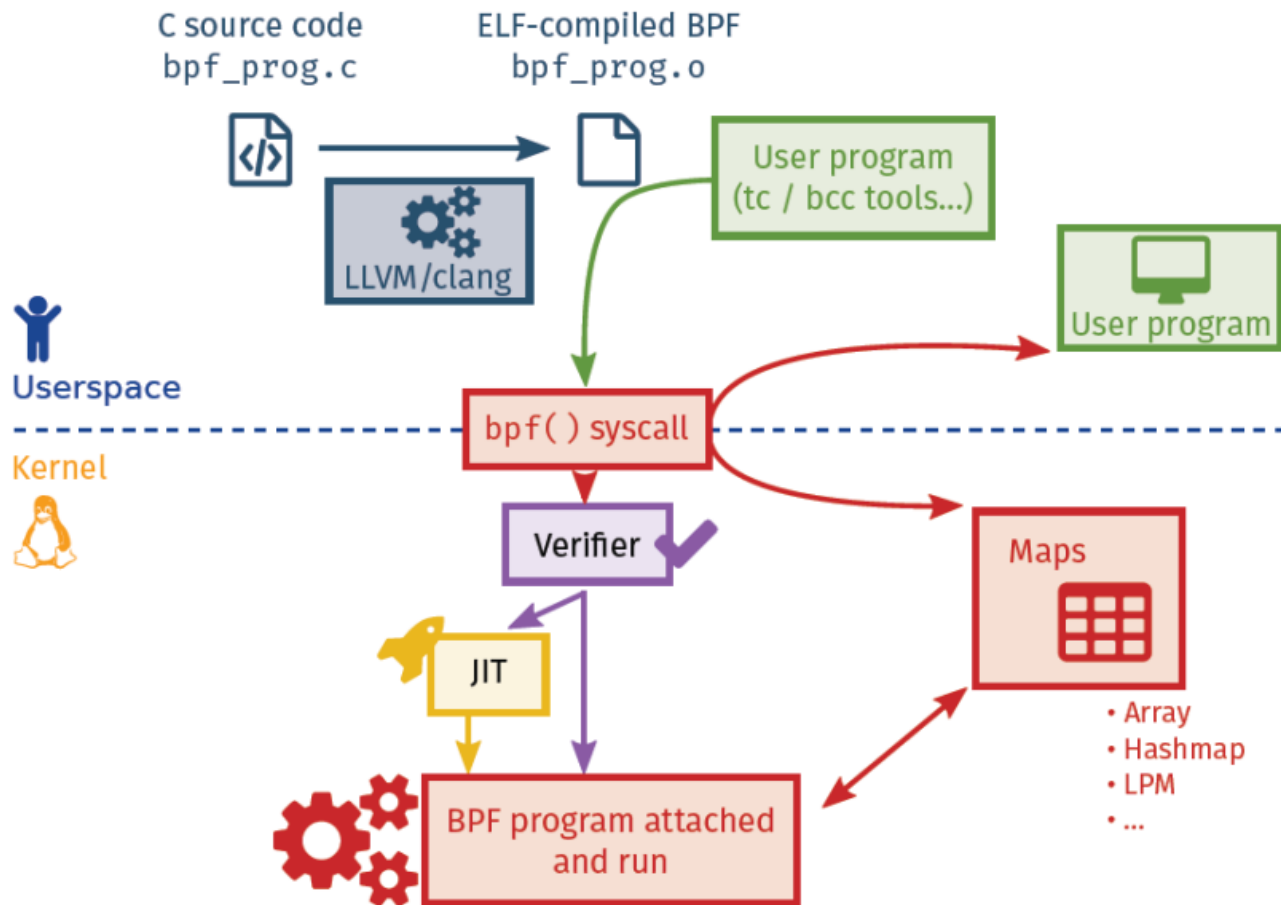


eBPF 框架



CHINA LINUX KERNEL
中国Linux内核开发者大会

- 加载 (load)
 - libbpf解析
 - 整合
 - 重定位
 - CO-RE
 -
 - verifier校验
 - 访问校验
 - 分支遍历
 - fixup改写
 -
 - JIT编译
 - BPF字节码 -> 机器码
- 挂载 (attach)
 - tracing
 - xdp
 -
- 触发 (trigger)
 - event
 -



RISC-V eBPF 实现



CHINA LINUX KERNEL
中国Linux内核开发者大会

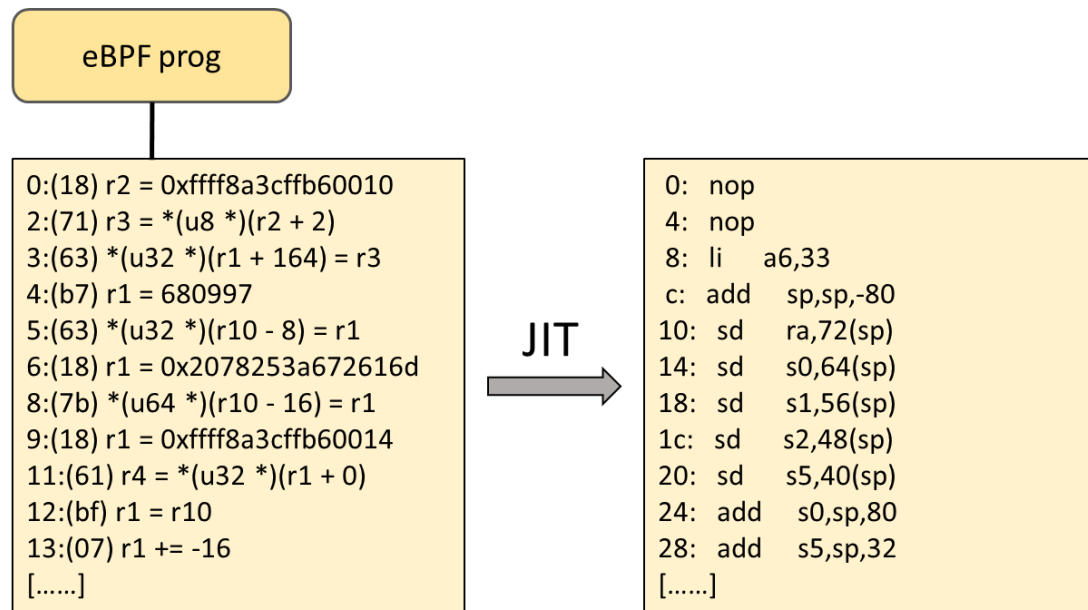
- JIT: 即时编译器, 将BPF字节码翻译成RV机器码

```
for (i = 0; i < NR_JIT_ITERATIONS; i++) { // 多轮JIT, 优化指令数
    emit prologue: 函数堆栈初始化, 保存寄存器等
    emit body: 将一条条BPF指令转换成RV机器码
    emit epilogue: 函数堆栈清理, 恢复寄存器等
}
```

- BPF trampoline

- 其他架构实现

- bpf_arch_text_poke
- bpf_arch_text_copy
- bpf_jit_alloc_exec
-





CHINA LINUX KERNEL
中国Linux内核开发者大会

02 RISC-V eBPF 现状

RV64 BPF JIT支持



CHINA LINUX KERNEL
中国Linux内核开发者大会

● 基本框架

- 寄存器映射: BPF寄存器<->RV寄存器

BPF寄存器	RISC-V寄存器
R0: return value	A5
R1-R5: arguments	A0-A4
R6-R9: callee saved	S1-S4
R10: FP	S5

- emit prologue
- emit body
- emit epilogue
- BPF基本指令

● 基本指令

- 运算: ADD/SUB/AND/OR等
- 跳转: JMP/CALL/TAILOCALL等
- 内存访问: LDX/STR等

```
author      Björn Töpel <bjorn.topel@gmail.com>      2019-02-05 13:41:22 +0100
committer   Daniel Borkmann <daniel@iogearbox.net>    2019-02-05 16:56:10 +0100
commit      2353ecc6f91fd15b893fa01bf85a1c7a823ee4f2 (patch)
tree        309a4f88c91f36813f16d88cc8627a737779b4fe /arch/riscv/net
parent      31de389707c8842ce71eaa8eff1eb74a43d5ef30 (diff)
download    linux-2353ecc6f91fd15b893fa01bf85a1c7a823ee4f2.tar.gz
```

bpf, riscv: add BPF JIT for RV64G

This commit adds a BPF JIT for RV64G.

The JIT is a two-pass JIT, and has a dynamic prolog/epilogue (similar to the MIPS64 BPF JIT) instead of static ones (e.g. x86_64).

At the moment the RISC-V Linux port does not support CONFIG_HAVE_KPROBES, which means that CONFIG_BPF_EVENTS is not supported. Thus, no tests involving BPF_PROG_TYPE_TRACEPOINT, BPF_PROG_TYPE_PERF_EVENT, BPF_PROG_TYPE_KPROBE and BPF_PROG_TYPE_RAW_TRACEPOINT passes.

The implementation does not support "far branching" (>4KiB).

Test results:

```
# modprobe test_bpf
test_bpf: Summary: 378 PASSED, 0 FAILED, [366/366 JIT'ed]
```

```
# echo 1 > /proc/sys/kernel/unprivileged_bpf_disabled
# ./test_verifier
```

...

```
Summary: 761 PASSED, 507 SKIPPED, 2 FAILED
```

RV64 BPF 压缩扩展支持



CHINA LINUX KERNEL
中国Linux内核开发者大会

● RISC-V指令集

- 模块化
- 通用指令集：I, M, A, F, D
- 扩展指令集：C, Zbb, Sscofpmf

● C Extension/压缩指令扩展

- 指令宽度：32-bit -> 16-bit
- 减少镜像体积
- 增加代码密度，提升性能

```
author      Luke Nelson <lukenels@cs.washington.edu> 2020-07-20 19:52:40 -0700
committer   Alexei Starovoitov <ast@kernel.org>    2020-07-21 13:26:25 -0700
commit      18a4d8c97b841632920c16a6fa9216d1214f3db7 (patch)
tree        08f67a054d33cdc161f9bf151de27a8b9ee7f763 /arch/riscv/net
parent      804ec72c68c8477b8713a1e8f8eda120d3471031 (diff)
download    linux-18a4d8c97b841632920c16a6fa9216d1214f3db7.tar.gz
```

bpf, riscv: Use compressed instructions in the rv64 JIT

This patch uses the RVC support and encodings from `bpf_jit.h` to optimize the rv64 jit.

The optimizations work by replacing `emit(rv_X(...))` with a call to a helper function `emit_X`, which will emit a compressed version of the instruction when possible, and when RVC is enabled.

The JIT continues to pass all tests in `lib/test_bpf.c`, and introduces no new failures to `test_verifier`; both with and without RVC being enabled.

Most changes are straightforward replacements of `emit(rv_X(...), ctx)` with `emit_X(..., ctx)`, with the following exceptions bearing mention;

* Change `emit_imm` to sign-extend the value in "lower", since the checks for RVC (and the instructions themselves) treat the value as signed. Otherwise, small negative immediates will not be recognized as encodable using an RVC instruction. For example, without this change, `emit_imm(rd, -1, ctx)` would cause lower to become 4095, which is not a 6b int even though a "c.li rd, -1" instruction suffices.

* For `{BPF_MOV, BPF_ADD}` `BPF_X`, drop using `addiw, addw` in the 32-bit cases since the values are zero-extended into the upper 32 bits in the following instructions anyways, and the addition commutes with zero-extension. (`BPF_SUB BPF_X` must still use `subw` since subtraction does not commute with zero-extension.)

RV64 BPF 异常处理表支持



CHINA LINUX KERNEL
中国Linux内核开发者大会

- 简化内存读取逻辑
 - `bpf_probe_read(dst, size, ptr_of_val)`
 - `dst = ptr_of_val`
- 保证内存读取安全
 - fixup when illegal memory access
 - 防止内核崩溃

```
author      Tong Tiangen <tongtiangen@huawei.com> 2021-10-27 11:18:22 +0000
committer   Daniel Borkmann <daniel@iogearbox.net> 2021-10-28 01:02:44 +0200
commit      252c765bd764a246a8bd516fabf6d6123df4a24f (patch)
tree        47ba312f3ec4698cd0829237b01d3f2bbc5ee492 /arch/riscv/net
parent      03e6a7a94001b9582ef6549e5709f3d684217b28 (diff)
download    linux-252c765bd764a246a8bd516fabf6d6123df4a24f.tar.gz
```

riscv, bpf: Add BPF exception tables

When a tracing BPF program attempts to read memory without using the `bpf_probe_read()` helper, the verifier marks the load instruction with the `BPF_PROBE_MEM` flag. Since the riscv JIT does not currently recognize this flag it falls back to the interpreter.

Add support for `BPF_PROBE_MEM`, by appending an exception table to the BPF program. If the load instruction causes a data abort, the fixup infrastructure finds the exception table and fixes up the fault, by clearing the destination register and jumping over the faulting instruction.

A more generic solution would add a "handler" field to the table entry, like on x86 and s390. The same issue in ARM64 is fixed in 800834285361 ("bpf, arm64: Add BPF exception tables").

RV64 BPF trampoline



CHINA LINUX KERNEL
中国Linux内核开发者大会

- 零开销Tracing

- 类似Kprobe/Kretprobe
- fentry, fexit

- 安全拦截

- LSM HOOK
- fmod_ret

- 可编程扩展

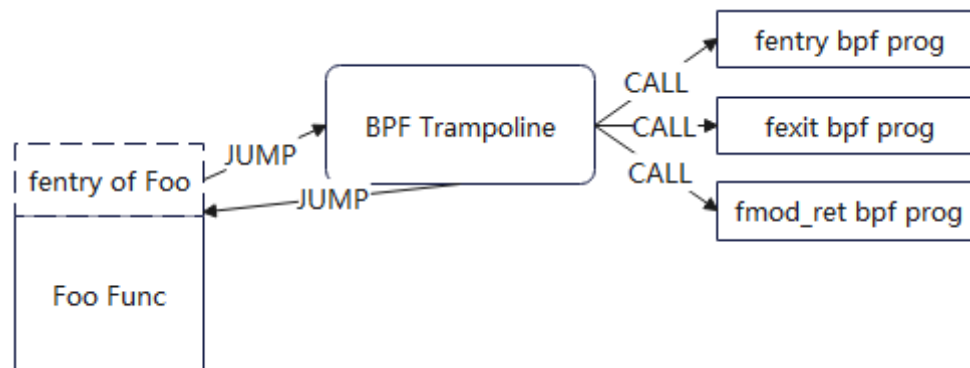
- 拥塞算法优化、sched_ext
- struct_ops

```
author      Pu Lehui <pulehui@huawei.com>      2023-02-15 21:52:05 +0800
committer   Daniel Borkmann <daniel@iogearbox.net> 2023-02-17 21:45:30 +0100
commit      49b5e77ae3e214acff4728595b4ac7bf776693ca (patch)
tree        8d8439f133914a38aa7aa4b7436726e9ca76ecba /arch/riscv/net
parent      596f2e6f9cf41436a5512a3f278c86da5c5598fb (diff)
download    linux-49b5e77ae3e214acff4728595b4ac7bf776693ca.tar.gz
```

riscv, bpf: Add bpf trampoline support for RV64

BPF trampoline is the critical infrastructure of the BPF subsystem, acting as a mediator between kernel functions and BPF programs. Numerous important features, such as using BPF program for zero overhead kernel introspection, rely on this key component. We can't wait to support bpf trampoline on RV64. The related tests have passed, as well as the test_verifier with no new failure cases.

```
Signed-off-by: Pu Lehui <pulehui@huawei.com>
Signed-off-by: Daniel Borkmann <daniel@iogearbox.net>
Tested-by: Björn Töpel <bjorn@rivosinc.com>
Acked-by: Björn Töpel <bjorn@rivosinc.com>
Link: https://lore.kernel.org/bpf/20230215135205.1411105-5-pulehui@huaweicloud.com
```



RV64 BPF kfunc



CHINA LINUX KERNEL
中国Linux内核开发者大会

- 支持调用kernel function
 - bpf helper: 新增复杂、ABI固定
 - 依赖BTF of kernel function
- 优点
 - 定制化 (kernel、module)
 - 易用性

author Pu Lehui <pulehui@huawei.com> 2023-02-21 22:06:56 +0800
committer Alexei Starovoitov <ast@kernel.org> 2023-02-22 13:11:31 -0800
commit d40c3847b485acc3522b62b020f77dcd38ca357f (patch)
tree 6864d578604c8050bfd72ef2394727b0834f6657 /arch/riscv/net
parent df2ccc180a2e6f6e4343eb99dcfab4f8af2816 (diff)
download linux-next-d40c3847b485acc3522b62b020f77dcd38ca357f.tar.gz

riscv, bpf: Add kfunc support for RV64

This patch adds kernel function call support for RV64. Since the offset from RV64 kernel and module functions to bpf programs is almost within the range of s32, the current infrastructure of RV64 is already sufficient for kfunc, so let's turn it on.

Suggested-by: Björn Töpel <bjorn@rivosinc.com>
Signed-off-by: Pu Lehui <pulehui@huawei.com>
Acked-by: Björn Töpel <bjorn@rivosinc.com>
Link: <https://lore.kernel.org/r/20230221140656.3480496-1-pulehui@huaweicloud.com>
Signed-off-by: Alexei Starovoitov <ast@kernel.org>

RV64 BPF pack allocator



CHINA LINUX KERNEL
中国Linux内核开发者大会

● 背景

- 系统存在大量bpf程序
- iTLB miss严重

● 实现

- 2M大页
- 聚拢小体积bpf程序

```
author      Puranjay Mohan <puranjay12@gmail.com> 2023-08-31 13:12:29 +0000
committer   Palmer Dabbelt <palmer@rivosinc.com> 2023-09-06 06:26:07 -0700
commit      48a8f78c50bd6f7f08fd40daa62252fd043f2f18 (patch)
tree        ad26a3315aca1638541a551bfa8c097ece5acced /arch/riscv/net
parent      cad539baa48ff257b598000a90db2b7edd4b2dd5 (diff)
download    linux-48a8f78c50bd6f7f08fd40daa62252fd043f2f18.tar.gz
```

bpf, riscv: use prog pack allocator in the BPF JIT

Use `bpf_jit_binary_pack_alloc()` for memory management of JIT binaries in RISC-V BPF JIT. The `bpf_jit_binary_pack_alloc` creates a pair of RW and RX buffers. The JIT writes the program into the RW buffer. When the JIT is done, the program is copied to the final RX buffer with `bpf_jit_binary_pack_finalize`.

Implement `bpf_arch_text_copy()` and `bpf_arch_text_invalidate()` for RISC-V JIT as these functions are required by `bpf_jit_binary_pack` allocator.

```
Signed-off-by: Puranjay Mohan <puranjay12@gmail.com>
Reviewed-by: Song Liu <song@kernel.org>
Reviewed-by: Pu Lehui <pulehui@huawei.com>
Acked-by: Björn Töpel <bjorn@kernel.org>
Tested-by: Björn Töpel <bjorn@rivosinc.com>
Acked-by: Daniel Borkmann <daniel@iogearbox.net>
Link: https://lore.kernel.org/r/20230831131229.497941-5-puranjay12@gmail.com
Signed-off-by: Palmer Dabbelt <palmer@rivosinc.com>
```


RV64 修复BPF和RISC-V ABI兼容



CHINA LINUX KERNEL
中国Linux内核开发者大会

● 背景

- BPF: 对32位imm进行零扩展
- riscv: 对32位imm进行符号扩展
- kfunc call: 影响
- helper call: 不影响, 总是u64

● 实现

- 对kfunc call的32-bit入参进行符号扩展

```
author      Pu Lehui <pulehui@huawei.com>      2024-03-24 10:33:06 +0000
committer   Alexei Starovoitov <ast@kernel.org> 2024-03-25 11:39:31 -0700
commit      443574b033876c85a35de4c65c14f7fe09222b2 (patch)
tree        6f8415231e5bb5407d6bc803e8f846c9d6379d04 /arch/riscv/net
parent      122fdbd2a030a95128737fc77e47df15a8f170c3 (diff)
download    linux-443574b033876c85a35de4c65c14f7fe09222b2.tar.gz
```

riscv, bpf: Fix kfunc parameters incompatibility between bpf and riscv abi

We encountered a failing case when running selftest in no_alu32 mode:

The failure case is `kfunc_call/kfunc_call_test4` and its source code is like bellow:

```
...
long bpf_kfunc_call_test4(signed char a, short b, int c, long d) __ksym;
int kfunc_call_test4(struct __sk_buff *skb)
{
    ...
    tmp = bpf_kfunc_call_test4(-3, -30, -200, -1000);
    ...
}
...
```

And its corresponding asm code is:

```
...
0: r1 = -3
1: r2 = -30
2: r3 = 0xffffffff38 # opcode: 18 03 00 00 38 ff ff ff 00 00 00 00 00 00 00
4: r4 = -1000
5: call bpf_kfunc_call_test4
...
```

insn 2 is parsed to ld_imm64 insn to emit 0x00000000ffffffff38 imm, and converted to int type and then send to bpf_kfunc_call_test4. But since it is zero-extended in the bpf calling convention, riscv jit will directly treat it as an unsigned 32-bit int value, and then fails with the message "actual 4294966063 != expected -1234".



CHINA LINUX KERNEL
中国Linux内核开发者大会

03 RISC-V eBPF 规划

RISC-V eBPF 规划



CHINA LINUX KERNEL
中国Linux内核开发者大会

- More Alignment
 - new features
 - exceptions, mixing tailcalls, etc.

- More Extensions
 - optimization, functionality
 - RVA22/RVA23

<https://github.com/riscv/riscv-profiles/blob/main/rva23-profile.adoc>


- More Robust
 - bugfix
 - refactor

- More Fast
 - insns count reduce
 - new algorithm/mechanism
 - rely on real hardware (VisionFive2)



CHINA LINUX KERNEL
中国Linux内核开发者大会

Thank you



群聊: CLK_2024_华为技术分享
群



该二维码7天内(10月28日前)有效, 重新进入将更新

APPENDIX



CHINA LINUX KERNEL
中国Linux内核开发者大会

- 一键RV64 BPF运行环境
<https://github.com/pulehui/riscv-bpf-vmtest>
- Specifications
<https://riscv.org/technical/specifications/>