



CHINA LINUX KERNEL
中国Linux内核开发者大会



华中科技大学
网络安全学院
School of Cyber Science and Engineering, HUST

第19届中国 Linux内核开发者大会



赞助单位



支持单位



支持社区&媒体



2024年10月 湖北·武汉



华中科技大学

vivo



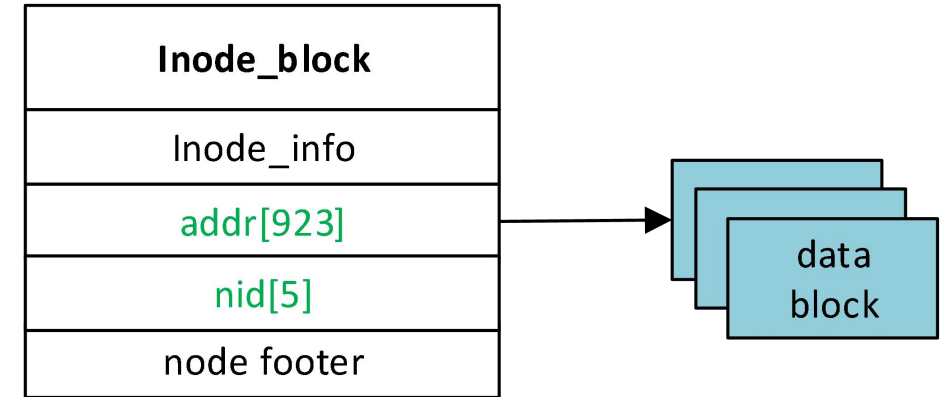

Inline Tail : 大幅减少 F2FS 小文件空间占用

伍勃 (bo.wu@vivo.com)

vivo 存储系统工程师

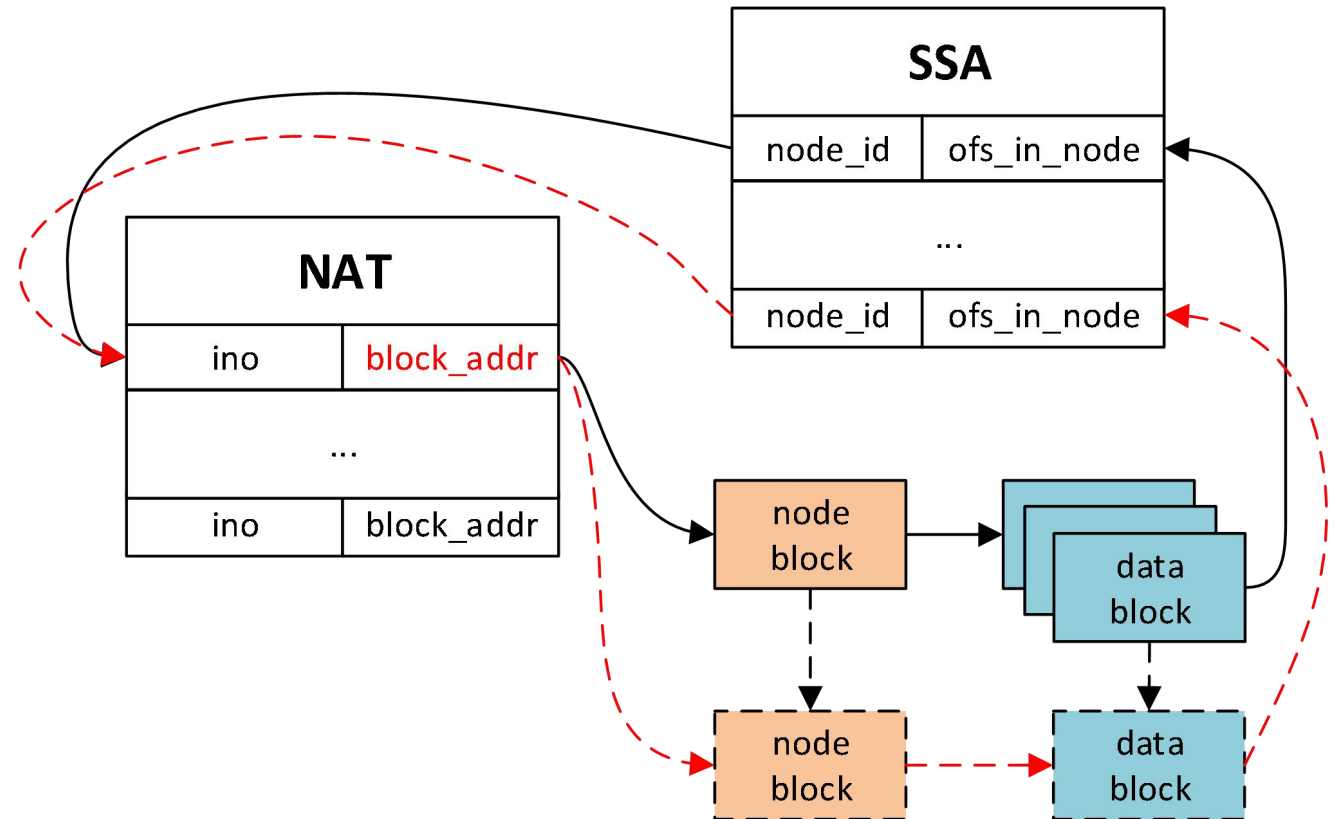
F2FS Inode is Too Large

- F2FS Inode occupies an entire 4kB block
- EXT4 : 256B, XFS : 512B
- File storage space = inode block + data block
- **Small files waste too much space**



Why is the F2FS Inode So Large

- F2FS defines two block types (node and data)
- Using NAT to avoid wandering tree problem
- **1 inode use 1 node block**
- Reducing inode size requires restructuring all metadata



F2FS Inode Space Utilization Optimization




- Xattr inline
 - Save xattr in inode
 - Default 200B
- Data inline
 - Save small file in inode
 - Size < 3.4kB
- **Small file size > 4kB?**

inode block	4096	xattr inline		data inline	
inode info	360				
addr table[923]	3692	extra info	0~36 (i_extra_isize)		
		addr[n]	3456~3492	reserved addr	4
				data inline	3452~3488
		inline xattr[50]	200		
nid table[5]	20				
node footer	24				

Principle of F2FS Inline Tail

- The tail block of a file is often not fully filled
- Address table of small files is essentially empty
- **Inline tail: store tail block into inode block**
- Can save one block(4kB)



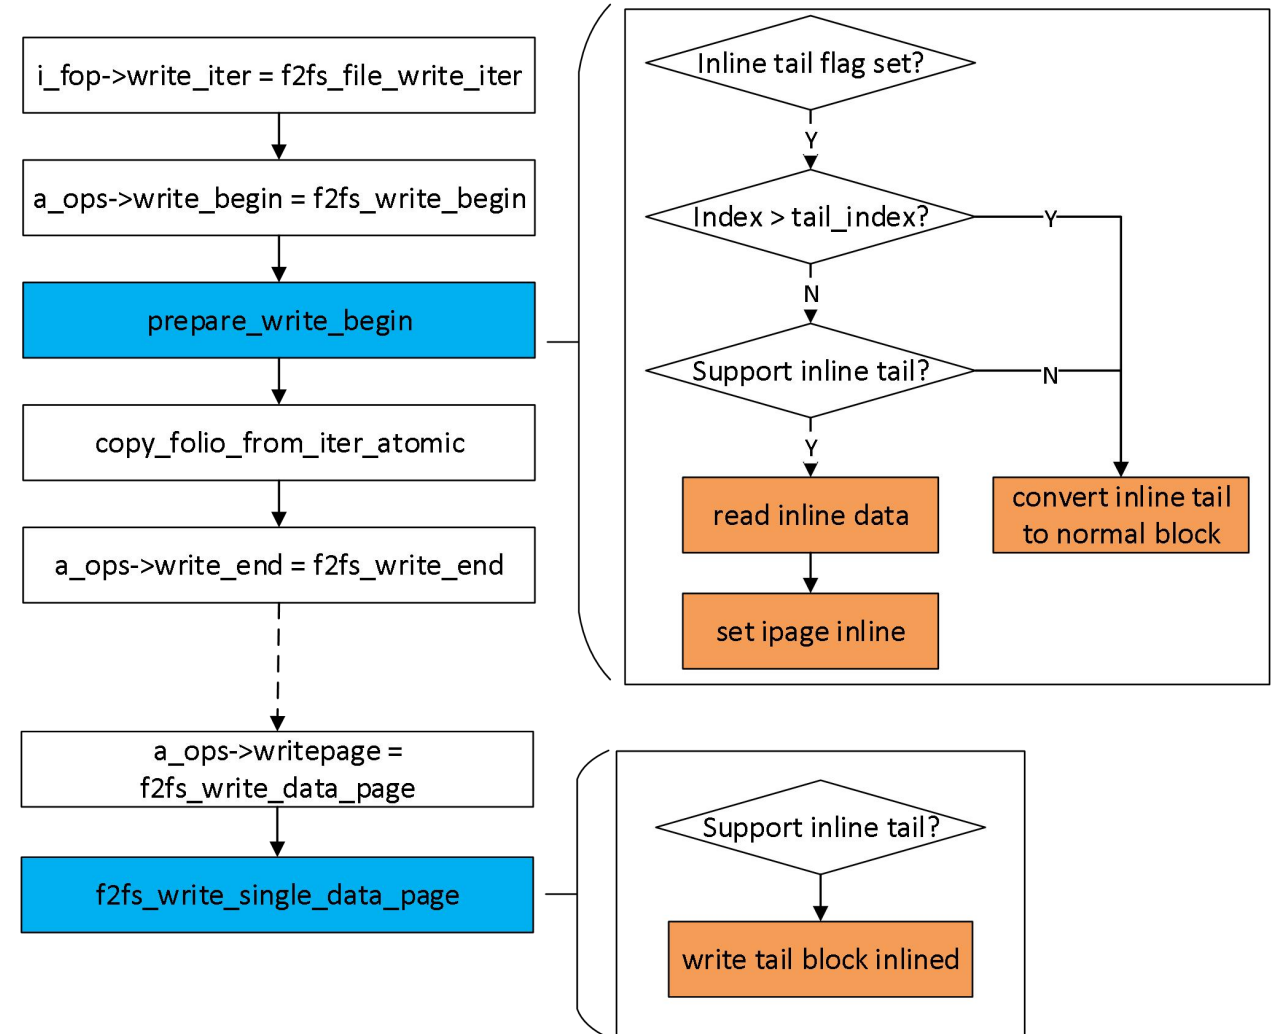
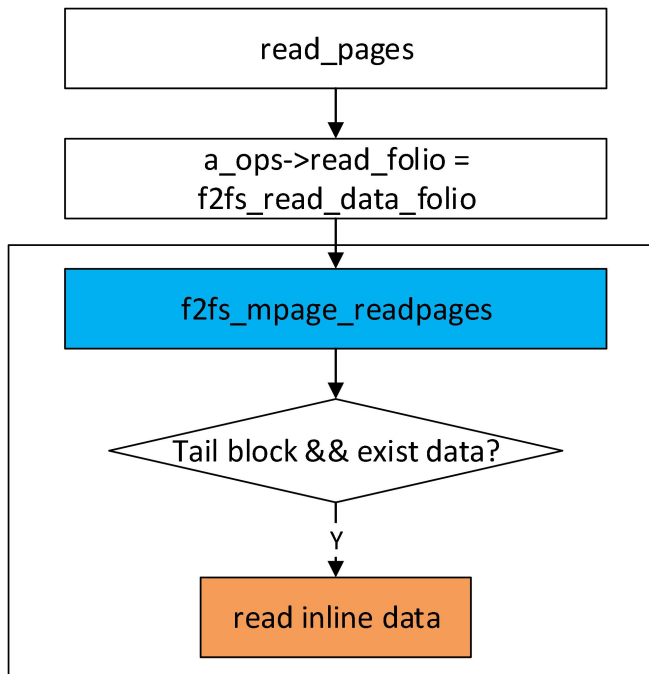
Design of F2FS Inline Tail

- **Size limitation: <68kB**
- Small files use a compact block address array with 16 entries
- Mixed blocks: tail block stored inline, others use traditional block
- Why 16 entries?: The marginal benefit decreases as the size increases

inode block	4096	tail inline	
inode info	360		
addr table[923]	3692	extra info	0~36 (i_extra_isize)
		addr[16]	64
		reserved addr	4
		inline tail	3392~3428
		inline xattr[50]	200
nid table[5]	20		
node footer	24		

Implementation of F2FS Inline Tail

- Only support buffer read & write
- Adapt inline tail for other FS interfaces



- Data race issue:
 - b3d208f96d6b ("f2fs: revisit inline_data to avoid data races and potential bugs")
 - Only support normal buffer write
 - **Inline conversion is permanently, not bidirectional**
 - Take lock when changing inline flag
 - Take lock when read file size
- Regression test:
 - Tool: xfstest-bld (<https://github.com/tytso/xfstests-bld>)
 - The test results are not very stable (random operations)

Benefits of F2FS Inline Tail

- Little files (4kB~68kB) can save one block(4kB), reduce most 1/3 space & IO request.
- Linux source code storage space -8%, copy time -10%
- Double the benefit from inline data

Inline Tail Saved Storage Space			
file size	w/o	/w	saved
<3.4k	4	4	0%
4k~7.4k	12	8	33%
8k~11.4k	16	12	25%
12k~15.4k	20	16	20%
16k~19.4k	24	20	17%
28k~31.4k	36	32	11%
44k~47.4k	52	48	8%
64k~67.4k	72	68	6%

- Replace inline data
 - **Inline data is a special case of inline tail**
- Measure the benefit of reducing IO request
 - 85% of files on user phones are smaller than 64kB
 - Reducing I/O should improve system performance
- Using inline tail as a file write buffer
 - Cache small synchronous data in the inline tail
- Support encryption on inline tail

Patches for F2FS Inline Tail

Wu Bo (13):

- f2fs: add inline tail mount option
- f2fs: add inline tail disk layout definition
- f2fs: implement inline tail write & truncate
- f2fs: implement inline tail read & fiemap
- f2fs: set inline tail flag when create inode
- f2fs: fix address info has been truncated
- f2fs: support seek for inline tail
- f2fs: convert inline tail when inode expand
- f2fs: fix data loss during inline tail writing
- f2fs: avoid inlining quota files
- f2fs: fix inline tail data lost
- f2fs: convert inline tails to avoid potential issues
- f2fs: implement inline tail forward recovery

fs/f2fs/data.c	93	+++++
fs/f2fs/f2fs.h	46	+++++
fs/f2fs/file.c	85	+++++
fs/f2fs/inline.c	159	+++++
fs/f2fs/inode.c	6	++
fs/f2fs/namei.c	3	+
fs/f2fs/node.c	6	+ -
fs/f2fs/recovery.c	9	++ -
fs/f2fs/super.c	25	+++++
fs/f2fs/verity.c	4	++
10 files changed, 409 insertions(+), 27 deletions(-)		

base-commit: 67784a74e258a467225f0e68335df77acd67b7ab



<https://lore.kernel.org/linux-f2fs-devel/cover.1726024116.git.bo.wu@vivo.com/>

THANK YOU